

DEPLOYMENT & OPERATIONS MANUAL

PROJECT: SYMBIOSE GEOSPATIAL TERMINAL

Next.js 15 — NestJS — PostGIS — Python LiDAR Engine

| | |
|---------------------|---|
| Component: | Full-Stack Geospatial Application |
| Frontend: | Next.js 15 (App Router) + Mapbox GL JS |
| Backend: | NestJS + TypeORM + Python 3 (Rasterio) |
| Database: | PostgreSQL 15 + PostGIS Extension |
| Target Env: | Docker Compose (Local / Prod) |
| Data Engine: | Stream-JSON ETL + Promise.all Spatial Query |

Contents

| | | |
|----------|--|----------|
| 1 | Architecture Overview | 2 |
| 2 | Data Preparation & Injection | 2 |
| 2.1 | Directory Structure Setup | 2 |
| 3 | Environment Configuration | 2 |
| 4 | Docker Compose Deployment (Local & VPS) | 3 |
| 4.1 | Build and Spin Up Services | 3 |
| 4.2 | Execute the ETL Seeding Pipeline | 3 |
| 4.3 | Access Points | 3 |
| 5 | Remote Production Roadmap | 3 |
| 6 | Troubleshooting & Optimization | 4 |

1 Architecture Overview

The Symbiose Geospatial Terminal is a high-performance, containerized application designed to process and visualize massive spatial datasets (BD Forêt, Etalab Cadastre, and IGN LiDAR HD).

- **Frontend:** Next.js serving a Mapbox GL JS map with dynamic viewport bounding-box guards (debounced) to prevent browser memory overload.
- **Backend:** NestJS providing RESTful endpoints. Utilizes parallel `Promise.all` execution to concurrently fetch PostGIS ST_Intersects data (Bio-Data & Cadastre) while spawning a Python child process for LiDAR masking.
- **ETL Pipeline:** A memory-efficient stream processing architecture (`stream-json`) engineered to handle $\geq 400k$ geo-features with a constant memory footprint ($\leq 100MB$).

2 Data Preparation & Injection

⚠ Critical Data Warning

Spatial datasets (Shapefiles, GeoJSON, GeoTIFF) scale up to several Gigabytes. **Never** commit the `data/` directory to version control. Download them manually to the host machine before building the containers.

2.1 Directory Structure Setup

Ensure the following exact structure exists at the project root:

```
mkdir -p data/BDV2 data/CPL data/lidar
```

1. **BD Forêt (Core):** Extract IGN shapefiles (.shp, .dbf, .prj) into `./data/BDV2/`.
2. **Cadastre (Bonus A):** Place Etalab Cadastre files (e.g., `cadastre-75-parcelles.json`) into `./data/CPL/`.
3. **LiDAR CHM (Bonus B):** Place IGNF MNH GeoTIFFs (.tif) into `./data/lidar/`.

3 Environment Configuration

Configure these variables in your `.env` file at the project root or inject them via your CI/CD pipeline:

| Variable Name | Description |
|---------------------------------|---|
| NEXT_PUBLIC_MAPBOX_ACCESS_TOKEN | Required. Mapbox API Token for rendering the base map. |
| DATABASE_HOST | Defaults to db in Docker Compose network. |
| JWT_SECRET | Secret key for securing API sessions and state persistence. |

4 Docker Compose Deployment (Local & VPS)

This project utilizes a Monorepo Docker orchestration. The backend container seamlessly integrates the Node.js runtime with Python/GDAL bindings.

4.1 Build and Spin Up Services

```
# Build images and start services in detached mode
docker-compose up --build -d
```

4.2 Execute the ETL Seeding Pipeline

Once the containers are healthy, you must ingest the raw files into PostGIS. This script handles EPSG:2154 to EPSG:4326 reprojection, encoding fixes, and spatial indexing.

```
# Trigger the high-performance stream pipeline
docker exec -it nest_backend pnpm run seed:forest
```

4.3 Access Points

- **Geospatial Terminal UI:** <http://localhost:3001>
- **Backend REST API (Swagger):** <http://localhost:3000/api>

5 Remote Production Roadmap

For a full production launch, the following architecture is recommended:

- **Frontend Hosting (Vercel):** Zero-config deployment. Ensures global CDN distribution for static assets and Mapbox tiles.

- **Backend App (Render / Railway):** Node.js Web Service. *Crucial:* Must attach a persistent Disk Volume to store LiDAR GeoTIFFs to bypass container ephemeral storage limits.
- **Database (Supabase / Render PostgreSQL):** Managed PostgreSQL instance with the PostGIS extension explicitly enabled (`CREATE EXTENSION postgis;`).

6 Troubleshooting & Optimization

| Symptom | Resolution |
|--------------------------------|--|
| OOM (Out of Memory) during ETL | If processing departments \geq 400k parcels (e.g., Dept 77), ensure the <code>stream-json</code> backpressure pipeline is active. Limit initial seeding to Dept 75 & 92 if host RAM is \leq 4GB. |
| Map Renders as Blank Grid | Invalid or missing <code>NEXT_PUBLIC_MAPBOX_ACCESS_TOKEN</code> . Rebuild the frontend container without cache. |
| No Cadastre Data (Bonus A) | Ensure zoom level is ≥ 13 . Verify the frontend payload is hitting the backend <code>ST_Intersects</code> concurrent logic. |

End of Document